



## TEACHING DATA SCIENCE TO STUDENTS IN BIOLOGY USING R, RSTUDIO AND LEARNR: ANALYSIS OF THREE YEARS DATA

GUYLIANN ENGELS <sup>\*1</sup>, PHILIPPE GROSJEAN <sup>1</sup> AND FREDERIQUE ARTUS <sup>2</sup>

<sup>1</sup> Numerical Ecology Department, Complexys and InforTech Institutes, University of Mons, Avenue du Champ de Mars, 8, 7000 Mons, Belgium

<sup>2</sup> Pedagogical Support and Quality Assurance Department, University of Mons, Place du Parc, 20, 700 Mons, Belgium

**ABSTRACT.** We examine the impact of implementing active pedagogical methodologies in three successive data science courses for a biology curriculum at the University of Mons, Belgium. Blended learning and flipped classroom approaches were adopted, with an emphasis on project-based biological data analysis. Four successive types of exercises of increasing difficulties were proposed to the students. Tutorials written with the R package learnr were identified as a critical step to transition between theory and the application of the concepts. The cognitive workload needed to complete the learnr tutorials was measured for the three courses and it was only lower for the last course, suggesting students needed a long time to get used to their software environment (R, RStudio and git). Data relative to students' activity, collected primarily from the ongoing assessment, were also used to establish student profiles according to their learning strategies. Several suboptimal strategies were observed and discussed. Finally, the timing of students contributions, and the intensity of teacher-learner interactions related to these contributions were analyzed before, during and after the mandatory distance learning due to the COVID-19 lockdown. A lag phase was visible at the beginning of the first lockdown, but the students' work was not markedly affected during the second lockdown period which lasted much longer.

**1. Introduction.** In a context where there is an exponentially growing mass of data [30], a reproducibility crisis in Science [4], and a progressive adoption of Open Science practices [5], statistics is broadened to a wider discipline called Data Science [13]. For the Data Science Association, “the Data Science means the scientific study of the creation, validation and transformation of data to create meaning” (<http://www.datascienceassn.org/code-of-conduct.html>). These changes have also led to the emergence of data science programs in universities and other higher education institutions [15, 12]. One example is the Harvard Data Science initiative (<https://datascience.harvard.edu/about>) launched in 2017. With a broader approach, also comes a broader audience. Such data science courses are not limited to computer scientists, mathematicians or statisticians. Students in humanities, social sciences, and natural sciences also attend them (for instance, the data science training at Duke University [12]). The focus of such courses is for students to develop the ability to deal with real datasets in all their complexities, to be able

2020 *Mathematics Subject Classification.* Primary: 58F15, 58F17; Secondary: 53C35.

*Key words and phrases.* Data Science Teaching, Flipped Classroom, R, RStudio, git.

\* Corresponding author: Guyliann Engels.

to conduct reproducible analyses, and to interpret these analyses in the light of knowledge in their field of expertise.

These data science courses pose several pedagogical challenges because numerous and unfamiliar concepts must be acquired by a heterogeneous class population [21]. Learning objectives span a large range of cognitive abilities and, in these courses, the intended learning outcomes aim to develop high-level cognitive process abilities, such as conceptual, procedural, and even metacognitive knowledge [26]. To meet such learning objectives, active learning methods are useful so that students can better acquire these high-level cognitive skills [18]. Advances in educational psychology, reviewed by Kirchner and Hendrick [24], give a scientific background to understand why a pedagogical practice does or does not work. Many teaching and learning frameworks involving numerical tools turn to a blended learning scenario, including remote activities to be done before and after in-class work, individual and group problem-solving, peer instruction and ongoing assessment. The flipped classroom approach also involves a mix between at home and in class activities, but learning occurs before the classes that are dedicated to discussions and problem solving. This allows students to be active in their learning, which has the benefit of improving student competences [18]. Moreover, this approach enables students to work at their own pace. Their diverse learning strategies are respected as they are actors in their own learning process [38]. Such frameworks are open learning centered and are supported by a varied and rich pedagogical environment [9].

The data transformation part of the process in data science is a challenge for students with little or no background at all in computing sciences. Students that do not master one or more computer languages enter an unfamiliar world and have to deal with many exotic concepts, techniques and tools. Version control systems like git, and their Internet hosting counterparts like GitHub, Gitlab or Bitbucket are also tools that are taught and used in data science courses [17, 23]. The use of document formats that dissociate content from presentation, namely LaTeX, Jupyter Notebook, R Markdown and Quarto to cite but a few, also contribute to the large number of potentially new tools learners have to discover [6]. On the other hand, a student in computing science already masters one or more computing languages. They are acquainted with version control systems, with databases and with the way data are manipulated and represented on a computer. Yet, the same students from computing science could have difficulties grasping the context of the study related to foreign disciplines. A student in mathematics or statistics is familiar with various concepts that underpin the techniques used to analyze the data. Students in biology, medicine, psychology, social sciences, economics, etc. obviously have very different *a priori* knowledge. The gap between knowledge and learning outcomes generates anxiety (see for instance [32]). The course must thus be organized in a way that learners progress little by little to avoid exposure to too many intimidating concepts and tools at once, taking into account their respective *a priori* knowledge and their initial gaps.

Suitable computer hardware and software environments are required to apply the concepts learned in data science courses. Different approaches range from using software accessed from a server [39] (RStudio Cloud (<https://rstudio.cloud/> [35]), Chromebook data science (<http://jhudatascience.org/chromebookdatascience/>)) to local installation on the student's computers. The former requires infrastructure to run the software on a server, and that software is only accessible to the students during the course. The latter raises problems of licensing for proprietary software

as well as installation and configuration issues. An intermediary solution uses pre-configured virtual machines, or containers (e.g. Docker) [11, 8]. Such a solution is the most flexible one because it can be deployed almost anywhere (in the computer lab, at home, on a laptop, etc.). As practical applications are important to learn data science [27], the correct choice of software is critical. Early exposure to the tools the students are most susceptible to use later in their work is desirable. This was highlighted by Auker and Barthelmess [3] for instance, for ecologists and by Alvarenga da Silva and Sampaio Moura [2] for physicians.

Recently, data science has also been used to analyze the effect of various pedagogical practices on the outcome of these courses thanks to learning analytics [16]. A vast amount of data can be collected about students' activities, and the analysis of these data allows the comparison of the impact of different pedagogical approaches, or to quantify and document the impact of such changes in the courses [33].

At the University of Mons in Belgium (UMONS), we reworked our biostatistics courses in the biology curriculum in 2018. A series of Data Science courses were introduced, both for our undergraduate and graduate students. The goal of these courses was to train biological data scientists capable of extracting meaningful information from raw biological data. They must be able to do so in a reproducible way and with the correct application of statistical tools and an appropriate critical mind. The goal was thus vastly broaden and it not longer solely targeted skills in biostatistics.

The learning outcomes defined for these Biological Data Science courses were thus to be able to analyze most recurrent biological data in practice and to present the results clearly and accurately in a scientific report. In order to achieve these learning outcomes the students had to master skills in biostatistics and scientific writing, and they had to become proficient in the use of computer tools like R, RStudio, git, GitHub and R Markdown. They also had to develop a critical mind in statistical thinking. All of these learning outcomes are described in the students' study programme at UMONS (see for instance for the academic year 2020-2021 [40, 41, 42]). A preconfigured VirtualBox machine with R, RStudio, R Markdown, git, and a series of preinstalled R packages was used (<https://www.sciviews.org/software/svbox/>) as a convenient means to deploy the same software environment both on the university computers and on the students' own laptops.

As our courses were reworked, we also decided to use flipped classroom and progressive adoption of suitable pedagogical practices in that particular context with a cyclical approach that consisted of stating goals, building pedagogical material with a large emphasis on numerical tools and collection of students' activities, and finally, analyzing the data collected. This approach allowed us to enhance our teaching activities the following academic year with improved pedagogical techniques. This approach is known as the educational data mining knowledge discovery cycle [33]. Here, we present the main results spanning three successive academic years from 2018-2019 to 2020-2021, including two particular periods where distance learning was forced due to the COVID-19 pandemic lockdown. In this paper, we will focus on the following three questions:

- Transition from theory to practice is critical and tutorials built with learnr (<https://rstudio.github.io/learnr/>) are capstones in our courses. What cognitive workload and perceived workload do these tutorials represent for students?

- How could we use learning analytics to spot suboptimal learning strategies and discriminate different student profiles in our biological data science courses?
- Did the quick shift from face-to-face to distance learning imposed by the COVID-19 lockdown periods affect our students' production and did it require increased exchanges with the teaching staff to support it?

**2. Methods.** This study focuses on three successive courses of increasing difficulty, referred to here as A, B and C and designed as a continuum. These courses were part of the core curriculum and were thus mandatory for all students enrolled on the Bachelor, or Master in Biology at UMONS.

- Course A was about data preparation, description and visualization. It also introduced inference with most common hypothesis tests in biology ( $t$  test, ANOVA, etc.). It was taught in the second year of the Bachelor. Course A was set up to assume only background knowledge acquired by the students during their first year of the same Bachelor in Biology at UMONS.
- Course B taught data modeling (linear, generalized linear and nonlinear models) and multivariate analyses (PCA, MDS, clustering, etc.) to students enrolled in the third year of the Bachelor in Biology at UMONS. Course A was a prerequisite to Course B (all students following Course B have previously passed Course A).
- Course C was taught to all Master students in the Biology section (first year of the Master). This course focused on machine learning, time series analysis and the analysis and visualization of georeferenced data. These students had either passed both Courses A and B at UMONS, or they demonstrated similar knowledge. Only a very few number of students (only one student in 2020-2021 out of a total of 26) came from a different Bachelor and thus had different courses, not including A and B, in their curriculum.

The course material was available online (<https://wp.sciviews.org>) and was centralized on a Wordpress site. Students had to login with their GitHub account and their academic data were collected from the UMONS Moodle server (<https://moodle.umons.ac.be>). The courses were broken down into modules that amounted of roughly 15h of work each. There were two in-class sessions of 2h and 4h per module (outside the lockdown periods, of course). There was roughly 3h of preparation at home before each session, and 3h of work to complete one module. The main activities in the class were analysis of actual data (projects). Students also asked questions and followed brief lectures (15 minutes) on selected topics in the class. They had to propose and vote for the topics to be covered during these short lectures. Finally, we encouraged students to help each other and to explain what they understood to their colleagues. Indeed, students' questions were sometimes redirected by the teacher to other students that had already mastered the topic. Teachers rarely answered questions directly. When it was possible, they rather proposed new tracks or ideas to investigate and helped learners to find the solution themselves. Students who went through the activities before the others were also encouraged to help their slower colleagues.

Regarding the timing, one module was taught every other week so that students had enough time to prepare the material at home before the in-class session, and after it, to finalize their projects for the module. As a term is made of 14 weeks, we did not teach more than six modules in a course unit to avoid teaching too much in a short time. After reading the theory, students were exposed to exercises of four

increasing levels of difficulty. They thus had to apply the concepts repeatedly but in different contexts, which broke any monotony and maintained a stimulating rhythm all along their progression. They had to learn the principles in the online book (<https://wp.sciviews.org>) and self-assessed their comprehension of the concepts using H5P (<https://h5p.org>) exercises (Level 1 difficulty). These exercises were simple questions (TRUE/FALSE, multiple choice, etc.). Learnr tutorials (<https://rstudio.github.io/learnr/>) were used for the Level 2 exercises. They were gently introduced the students to the R code required for the analyses and guided them step by step through their first data analysis. These tutorials were thus the entry point to the practice.

Most of the practical work was dedicated to GitHub projects (Level 3 and Level 4 difficulties). At this stage, the use of R instructions was not sufficient to complete the exercises. Students had to also become acquainted with git, GitHub, R Markdown and RStudio to manage the projects. They also had to interpret the results they obtained. The individual projects (Level 3) contained guidance on how to perform the different steps of the analyses. The group projects (Level 4, groups of two or four students) did not contain such guidance. At Level 3, the goals were clearly specified in the projects. At Level 4, students had to imagine suitable biological and statistical questions that could be answered by analyzing the data proposed in the project. Working on these projects represented both the core of in-class activities and the best expression of their learning progression. By construction, Level 1 to Level 4 exercises were built according to their increasing cognitive difficulties following bloom's taxonomy [26].

All student activities in H5P exercises (self-assessing), and in the learnr tutorials (transitioning smoothly from theory to practice) were recorded in a MongoDB database. The {learnitdown} R package (<https://www.sciviews.org/learnitdown/>) provided the code required to manage user login, user identification and activity tracking for this interactive material.

Projects containing the data, the analyses and the reports were hosted in GitHub repositories. These repositories were cloned and edited by the students in their virtual machines (SciViews Box) with RStudio (<https://www.rstudio.com/products/rstudio/>), either on their laptops or on the computers in the lab. We encouraged our students to install the virtual machine for the course on their own computer so that they were able to work comfortably at home and could also use it for other activities too. Assignment and creation of the GitHub repositories for each student, or group of students, was orchestrated by GitHub Classroom (<https://classroom.github.com>). Reports were written in R Markdown (<https://rmarkdown.rstudio.com/>), a file format that combines the prose with R code to produce analysis results, plots and tables directly inside the documents. All repositories were ultimately cloned by the teacher in a centralized area on our servers and data about commits (git logs) were collected using git version 2.31.1 and R version 4.0.5 [34]. To give an idea of the data recorded in 2020-2021, we had just over 3,500 events that were recorded for each student.

In distance learning, student support was done via email and Discord (<https://discord.com>). At the end of an academic term, all recorded messages were collected into text files. These files were scraped using custom R code to create a table with key information (basically, who, when, and what) for each message. Surveys were done periodically in class through Wooclap questionnaires (<https://www.wooclap.com>). Such questionnaires were used to query perceived workload

TABLE 1. Four levels of increasing difficulties in the exercises.

Level	Description	Type
L1	Interactive exercise in the course, direct feedback	h5p
L2	Tutorial with guided exercises, feedback and hints	learnr
L3	Individual and guided data analysis	individual project
L4	Free data analysis and reporting (by 2 or 4 students)	group project

of the learnr tutorials. Results were manually exported out of Woodclap by means of Excel files. These data were then incorporated into a table in our database thanks to an R script.

Data about users, courses, lectures and projects, as well as grading items (on average, more than 130 grading items were established for each student in 2020-2021) were pseudonymized: names, emails and all the personal information were replaced by random identifiers. The different tables were ultimately exported into CSV files and made public [19]. Data collection, treatment and use respect the European GDPR (General Data Protection Regulation) since each student had to agree explicitly with the way data were collected and used (including the research purpose) before each course started. They were able to visualize their data through personalized reports at any time.

The course material was organized in a way that favored autonomy and self-assessment (direct feedback in the exercises, hints and retry buttons in case of wrong answers). Table 1 summarizes the main characteristics of the exercises according to the difficulty level.

R and tidyverse [44] packages (<https://www.tidyverse.org>) were used to prepare the data and the analyses. The pseudonymized data are available from Zenodo [19, 20]. A GitHub repository with the code used to create the figures and table in this paper is available at [https://github.com/BioDataScience-Course/teaching\\_data\\_science\\_in\\_biology](https://github.com/BioDataScience-Course/teaching_data_science_in_biology).

### 2.1. Measured and Perceived Cognitive Workload in learnr Tutorials.

The average number of trials that were required for each student to find the right answer in learnr tutorial exercises was used as a proxy of measured workload. In comparison, the perceived cognitive workload was established with a NASA LTX questionnaire. This questionnaire is composed of six questions on a Likert scale [22]. The questions concern mental load, physical load, time pressure, expected success, effort required, and frustration experienced during the accomplishment of the task. The average value for the six questions constitutes a Raw Task Load index (RTLX) [10] that we used to quantify how students felt when using these learnr tutorials. An analysis of variance test and a Tukey's post-hoc Honest Significant Difference (HSD) tests were used for the comparison between courses.

### 2.2. Students Activity Profiles with Ongoing Assessment.

Data from the L1-L4 exercises and the student support (email and Discord) were used to characterize the students' activity profiles. A non-supervised classification technique called a Self-Organizing Map (SOM) [25] was used to characterize various learning profiles. The {kohonen} R package was used to compute the model [43].

### 2.3. Transition from Face-to-face to Distance Learning Imposed by the COVID-19 Lockdown.

The transition between face-to-face and distance learning



TABLE 2. Number of students, modules, and exercises for each course. For the learnr tutorials, the first number is the amount of tutorial documents and the second number in brackets is the total number of questions in these tutorials (year 2020-2021).

Course	Students	Modules	H5P	Learnr	Indiv. projects	Group projects
A	59	12	59	24 (211)	10	4
B	45	8	29	11 (108)	12	2
C	26	6	19	7 (37)	7	1

was studied through the contribution of each student to projects with the commits (git logs) and by these contributions were then divided by the questions they asked by email and on Discord. Though only descriptive analysis of these data was done, interesting patterns were observed.

**3. Results.** This study was performed on data related to the three successive courses that comprised 26 modules in total in 2020-2021. Table 2 summarizes the number of H5P, learnr, individual and group GitHub projects that students had to complete. Group projects usually spanned over several modules. It should be noted that for Course C, we also introduced a challenge in machine learning that replaced one GitHub group project. This challenge is omitted from the present analysis, being a unique activity that is difficult to compare to the rest. However, this explains why there was only one group project in Course C.

Retrospective data 2019-2020 [20] were also used when pertinent. It should be kept in mind that the pedagogical material was written and improved progressively over the three academic years. The H5P exercises and the auto-checking of learnr answers were not available before 2020-2021.

**3.1. Measured and Perceived Cognitive Workload in learnr Tutorials.** In our courses, learnr tutorials helped students to transition from the theory (online book chapters) to practice (projects). These tutorials are online interactive documents that recall main concepts, and take the students by the hand to perform their first data analysis step by step. At each step, they have at least one exercise or one quiz. The exercise consists of writing R code, or filling missing parts in R code to progress through the analysis.

Our goal with these tutorials was to optimally prepare the students for the practice of data science. The usefulness of these tutorials was qualitatively determined by observing the behavior of the students when they started their practical work. The number of retries necessary to complete an exercise on average, the number of exercises correctly answered, or the time needed to complete one tutorial are quantitative measurements and could be analyzed in order to optimize these tutorials.

A few tutorials were elaborated during the academic year 2018-2019, and positive feedback on their utility (both from direct observation of the students and thanks to their remarks) led us to systematize them into what we now call Level 2 activities (see Table 2) in the form of learnr documents in 2019-2020. The tutorials were further refined in 2020-2021: we added contextual hints with the {gradethis} R package (<https://pkgs.rstudio.com/gradethis/>). In their latest version, when students submit their answer to the exercises, the R code is parsed, analyzed and the result is compared with the solution. In case of differences, heuristics are used




FIGURE 1. Logarithm of the average number of retries that were required for each student to find the right answer in learnr tutorials exercises (2020-2021 academic year). This measure is used as a proxy to quantify the cognitive workload (with caution as explained in the text). The big black diamond is the average for the whole class. The small black dots are the outliers of the boxplot. The gray dots are the actual observations with small random horizontal shift to separate them. The number  $n$  above each box is the number of students.

to provide contextual hints. Students can then refine their solution and resubmit it. This appears very efficient in self-learning and self-assessing their competences before switching to the practice with confidence.

A fully objective quantitative measurement of the cognitive workload is near impossible to obtain in these asynchronous activities done at home. It was thus estimated by using a proxy: the logarithm of the average number of trials that were required for each student to find the right answer in the learnr tutorial exercises (Fig. 1). Caution is required here as a high number of trials could also be the result of students that are just guessing. However, answers being pieces of R code, pure guessing most probably leads to nothing useful. A certain level of understanding of both the R syntax and the question are required to obtain correct answers. Only data from students that correctly answered most of the exercises ( $> 90\%$ ) were used here, as it also rules out the students that appeared to have insufficient knowledge to master the concepts in the tutorials and were probably just guessing. For Course A, 12 out of 59 students did not pass this filter, 5 out of 45 for Course B and 1 out of 26 for Course C.  $\text{Log}(\text{trials} / \text{learnr exercise})$  is (mean  $\pm$  standard deviation):  $1.38 \pm 0.26$  for Course A,  $1.38 \pm 0.24$  for Course B and  $1.21 \pm 0.25$  for






FIGURE 2. Perceived workload for the learnr tutorials in the three courses (year 2020-2021). The big black diamond is the mean RTLX value. The small black dots are the outliers of the boxplot. The gray dots are the actual observations with small random horizontal shift to separate them. The number n above each box is the number of anonymous respondents.

Course C. At least one course differs significantly at 5% alpha level from the other two (ANOVA,  $F(2, 109) = 4.49$ ,  $p\text{-value} = 0.013$ ). The homogeneity of variances (Bartlett Test,  $K^2 = 0.30$ ,  $df = 2$ ,  $p\text{-value} = 0.86$ ) and the Normal distribution of the residuals using a quantile-quantile plot were verified. The students on Course C need significantly fewer trials to find the right answer than students on Courses A at  $\alpha$  level of 5% (Tukey HSD,  $t = -2.74$ ,  $p\text{-value} = 0.019$ ) and B (Tukey HSD,  $t = -2.68$ ,  $p\text{-value} = 0.023$ ).

The perceived cognitive load required to perform these exercises was also determined on the same students and for the same exercises. The Raw Task Load indexX (RTLX) measured the emotional state of the students after having completed a tutorial. This has, as far as we know, not yet been done. We used a NASA LTX questionnaire to assess it across all three courses. Participation in the survey was high: 48/59 (81%), 35/45 (78%) and 18/26 (69%) for Courses A, B, and C respectively.

The RTLX are (mean  $\pm$  standard deviation)  $52.1 \pm 9.7$  for Course A,  $49.8 \pm 11.9$  for Course B and  $43.7 \pm 14.2$  for Course C. The difficulty of the course, and thus, of the exercises in the tutorials increased from one course to the other. However, we did not observe any increase, neither in the number of retries, nor in the RTLX (Fig. 2). On the contrary, these appeared significantly lower for Course C than for Course A at the  $\alpha$  level of 5% (ANOVA,  $F(2, 98) = 3.59$ ,  $p\text{-value} =$

0.031; Tukey HSD,  $t = -2.68$ ,  $p\text{-value} = 0.023$ ). The homogeneity of variances (Bartlett Test,  $K^2 = 4.17$ ,  $df = 2$ ,  $p\text{-value} = 0.12$ ) and the distribution of the residuals using a quantile-quantile plot were verified and indicated that they did not depart significantly from a Normal distribution. The cognitive load perceived by the students diminished at the same pace as their ability to find the right answer in fewer trials. This may be a consequence of a more fluent R coding ability and a better mastery of the software environment.

**3.2. Student’ Activity Profiles with Ongoing Assessment.** The flipped classroom approach and the proactive attitude we expected from our learners (they had to formulate questions correctly whenever they faced a problem) led to different and contrasted learning strategies. Not all students asked questions. Some of them tried to find solutions on their own. Others preferred to ask their questions privately, while others had no problems exposing their difficulties in a public forum (the Discord channel dedicated to the course). The way and the timing learners progressed in the exercises also varied largely. The schedule was not tight and only suggested a rate of progression. No students were penalized if the exercises were done later, as long as they were completed before the deadline. As expected, a part of our students preferred to stick to the proposed schedule, while others procrastinated and delayed the completion of their exercises. Some strategies are probably more efficient than others. We analyzed the records of the students’ activities to distinguish the different learning profiles and we compared them with the grade they obtained at the end of the course.

In 2020-2021, to support the ongoing assessment without a final exam, the activity of each student in Level 1 (H5P) and 2 (learnr) exercises was exhaustively recorded in a database. For the GitHub projects (Levels 3 and 4 exercises), the GitHub repositories and the git log data were analyzed. During the lockdown periods, exchanges with students and answers to their questions were exclusively done by email, text or voice messages on Discord, or on private or public channels. Students were allowed to freely choose their favorite tool to interact with the teachers and between each other. All these exchanges were recorded too.

The degree of completion of all the exercises was used to establish the final grade for the course, with a much higher weight on individual and, especially, on group projects. The weight was adjusted from course to course according to the importance of the different projects. To give an idea, for Course A during the second term, Level 1 H5P exercises accounted for 5%, 10% for Level 2 learnr tutorials, 35% for Level 3 individual projects and 50% for Level 4 group works. On average, each student received more than 130 assessment items that accounted for their final grade. Two thirds of these assessments were established manually, using evaluation grids based on their work in the various projects. The remaining third was made of scores automatically calculated from the online exercises.

For the three courses, we recorded more than 450,000 events, which makes on average almost 3,500 events for each student. These data contain information that we used to characterize the behavior and learning patterns that the students exhibited. They are summarized into sixteen metrics.

For H5P exercises:

- trials/H5P ex.: the average number of trials for each H5P exercise until the right answer is found (students can retry as many times as they wish and they have immediate feedback on whether or not their answer is correct),

- correct H5P ex.: the fraction of H5P exercises that were answered correctly.

For learnr tutorial exercises:

- trials/learnr ex.: the average number of trials for each learnr exercise until the right answer is found (here also, students can retry as many times as they want), excluding quizzes,
- hints/learnr ex.: in learnr exercises, students can display hints to help them solve the problems (but they lose 10% of the exercise score for each hint they reveal). This is the average number of hints per exercise that were displayed for each student,
- correct learnr ex.: the fraction of learnr exercises that were completed with a correct answer,
- time/learnr ex.: the average time required to finish one learnr exercise involving R code writing, thus excluding quizzes.

For individual and group projects:

- commits/ind. projects: the average number of commits done by a student in one individual project,
- contributions/ind. projects: the number of lines changed -added or subtracted- in the R Markdown reports by one student in one individual project (this includes embedded R code for the processing, analysis and plotting of data),
- commits/group projects: same as above, but for group projects,
- contributions/group projects: same as above, but for group projects,
- percentage of contributions to group projects: the fraction of work the student did, relative to all the work done in group projects.

For support:

- questions/module: the number of questions students asked, divided by the number of modules in the course,
- percent of public questions: the fraction of questions that the student posted in a public channel (the Discord channel dedicated to the course that all the other students of the class can read),
- contributions/question: a metric that catches the relative “productivity” of the student related to the number of questions they ask.

Finally, global measurements:

- work done: the fraction of all exercises that the student completed,
- work done in time: the fraction of exercises done in the right time, that is, within the proposed schedule.

In our courses, we have a few students in mobility that come from various origins. The *a priori* knowledge is important in education. So, to avoid biases due to the past curriculum of the students, we restricted this analysis to the subpopulation that comes from the first year of the Bachelor in Biology at UMONS only. A Kohonen’s self-organizing map was used to create student profiles according to their activities (Fig. 3). A three-by-three hexagonal map was chosen, and students were thus classified into nine different groups.

In Fig. 3, the small peripheral plots in gray scale show how selected metrics distribute in the nine cells, from lowest value in white to highest value in black. They help to decipher the way students behave according to their profile. Metrics that are not represented in the figure exhibit similar patterns to others (for instance H5P metrics have a similar pattern to learnr metrics). The codebook vectors of

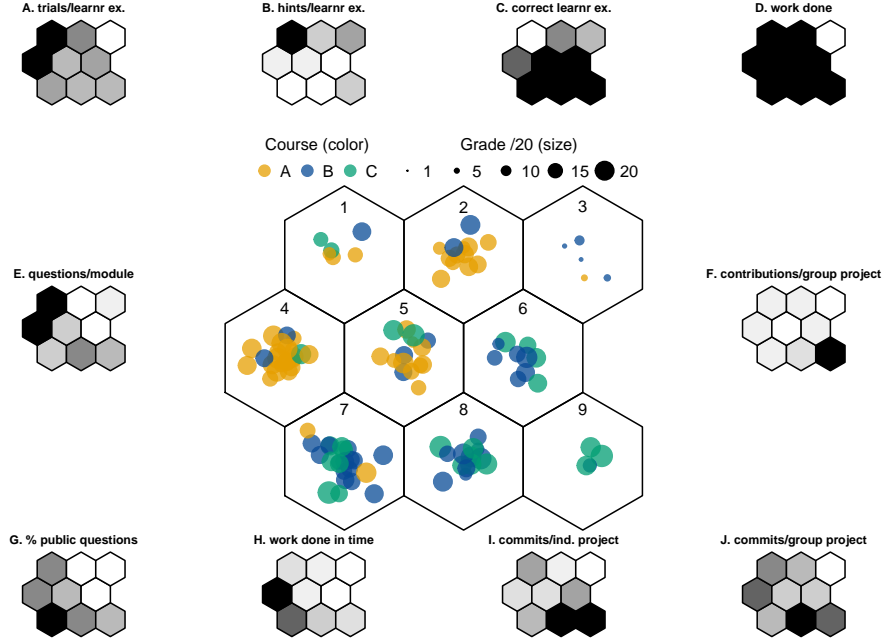


FIGURE 3. Self-organizing map of the student activities across the three courses (year 2020-2021). See the text for the explanations.

the self-organizing map is presented in the Appendix. These codebook vectors represent the importance of each variable in each cell. Dots in the central plot are the various students, with the color representing the course and the diameter of the dots indicating the grade the students obtained at the end of the course. The following paragraphs detail the information in this figure. The numbers between parentheses are the cell numbers in the central plot, and the uppercase letters in parentheses refer to the peripheral subplots.

Each cell (1-9) represents a learning strategy used by the students. The metric used to determine the effectiveness of the learning strategy is the grade each student obtained for the course. Learning strategies associated with high grade are qualified as “optimal”, and those corresponding to low grades are qualified as “suboptimal”. No causal relationship between the two is inferred: good or bad grades could result from other variables not studied here. We observe that suboptimal strategies are also associated with the time students need to learn (e.g., large number of trials per exercise) or to a lack of interaction in the group (e.g., very few questions asked).

Although most students completed all, or almost all of the exercises (D), Cell (3) collects the few students that did only a tiny part of these exercises. These students obtained very low grades. They belong to Courses A and B. On the other hand, heavy workers are at the bottom (I & J), and good performers in learnr tutorials (C) are in Cells (5-9).

- Cells (2) and (6) collect students that seldom asked questions (E), and that rarely appeared on the public channel (G). Minor differences separate them. For instance, learners in Cell (2) sometimes used hints (B), while those in

Cell (6) never did so, also because they found the correct answer to the exercises more often by themselves (C). Asking questions is at the core of our pedagogical approach. So, these students did not play the game. However, they were possibly successful anyway. Some of them probably exchanged with other students through different channels that we did not monitor. Cell (2) -more difficulty with learnr tutorials- mainly contains students belonging to Course A, while cell (6) contains students of Courses B and C. There is a clear evolution in their behavior from one course to the other in terms of ease in carrying out the exercises, even though they remained silent at the teacher-learner interaction level.

- Among the students that had a hard time figuring out the answers to the auto-evaluation exercises, Cell (1) reassembles people that most heavily relied on hints (B), and were among those who had to retry the exercises more often before they worked out the correct answer (A), a characteristic they share with Cell (4). These students also asked a lot of questions (E), both on the public and private channels (G, mid gray indicating a balance between public and private messages). The main difference between these two groups is that students in Cell (4) tried harder to find the answer without looking at the hints, while in Cell (1) they gave up more rapidly. Also these students respected the proposed schedule much more closely than all others (H). We have students coming from all courses there, but a majority from Course A.
- Cells 1-4 plus 6 contain students that exhibit suboptimal behaviors in one or the other way. The remaining Cells (5, 7-9) correspond to learner profiles that perform better from this point of view. Cell (5) is primarily represented by students from course A, but secondly, also from Courses B and C. These are average actors in all metrics, except they are fluent with Level 1 (H5P, not shown) and Level 2 (C) exercises.
- Moving from Cell (5) to (7), (8) and (9), we encounter increasingly good performers. The number of students from Course A becomes progressively lower, while Course B and, especially C, dominate in these groups. In Cell (7), they intensively used the public channel (G) and also respected the schedule quite well (H) as main differences from those from Cell (5). Students in Cells (8) and (9) were not so often in time, but this is because they were heavier workers in the projects, both in the individual (I) and in the group (J) activities. This obviously needed more time. In cell (9) we also find the students that contributed the most to the reports in terms of lines added or deleted (F).

To summarize, at the top of the SOM map, Cells (1-4, plus 6) contain students with suboptimal behaviors, Cell (5) are average students, and Cells (7-9) at the bottom exhibit profiles corresponding to the best performers. The pattern is also visible between Courses A (mainly distributed at the top or center of the map) to B and C (more represented at the bottom). This probably suggests that students needed time to get used to the course, its pedagogical approach, and/or the software environment they had to use. Since only a small fraction of the participating students failed, excluding the failing ones in Cell (4), the intercourse pattern can hardly be explained by a filter that eliminated the low performers from one course to the other.

**3.3. Transition from Face-to-face to Distance Learning Imposed by the COVID-19 Lockdown.** Due to the COVID-19 lockdown periods, distance learning had to be adopted abruptly. We analyzed the activity collected during the 2019-2020 and 2020-2021 academic years to describe the impact of these transitions on the progression of the students. In Fig. 4, the academic term is divided into seven work periods of approximately two weeks each (note that this was the suggested rate of the courses: one module every other week). The classes of the second term started in period Y1P09, since period Y1P08 was reserved for the first term exams session. The courses of the first term of 2020-2021 began in Y2P01. First lockdown started in period Y1P11 for one month and a half. Second lockdown started in Y2P03 and lasted to the end of the second term (Y2P15). During the first lockdown, we quickly opened the dedicated Discord channels that were available without any latency.

Contributions per student (Fig. 4a) was relatively constant during the second year, starting essentially in Y2P03, when the second lockdown was established. The highest activity is observable at the end (Y2P15), although there was no module taught during that period. This is because of the late students that finalized their reports at the last minute. Y2P01, Y2P02 and Y2P09 exhibit the lowest activity, and these were the start of the first and second terms. Y2P01 and Y2P01 were also taught in face-to-face and they correspond to the start of all three courses.

The learners-teachers interactions, especially during the lockdown, are quantified here by the number of questions. Figure 4b) shows the amount of contributions divided by the number of questions as a measure of work that was done on average by the students for each interaction. A higher value means more autonomy. A lower value indicates more problems or difficulties that require learners-teacher interactions to be solved. As this measure spreads over several orders of magnitude from one student to the other, a logarithmic scale is used. However, the median value -the bar inside the boxes- varies much less. the global number of questions during each period is less variable, as are the absolute contributions, leading to a rather stable ratio. The highest median ratios are observed at the last period of each term (Y2P07 and Y2P15) although no module was taught during that time. More contributions are observed relative to the questions at the end: students essentially finalize their reports.

The first year shows a different pattern. First, the lockdown period was restricted to the very end of the second term. Only the last module in both Courses A and B remained. In Y1P12, when distance learning was first imposed, we observed a marked decrease in the contributions per student (Fig. 4a). It is heavily compensated, and even overcompensated, in periods Y1P13 and Y1P14, which were by far the busiest periods of all. Period Y1P15 was not represented because it is after the deadline to finish all work that year.

The intensity of support during the first year shows a similar pattern as for the second year: extremely widespread from student to student. The median value is similar too, if not among the highest during periods Y1P11, Y1P13 and Y1P14. The productivity was thus not affected during that first lockdown, after a short lag time observable in Y1P12.

**4. Discussion.** Teaching data science to a population of students that are not very used to advanced computer techniques and tools, and that have only basic knowledge in mathematics and statistics is a hard task [37]. In order to let them learn

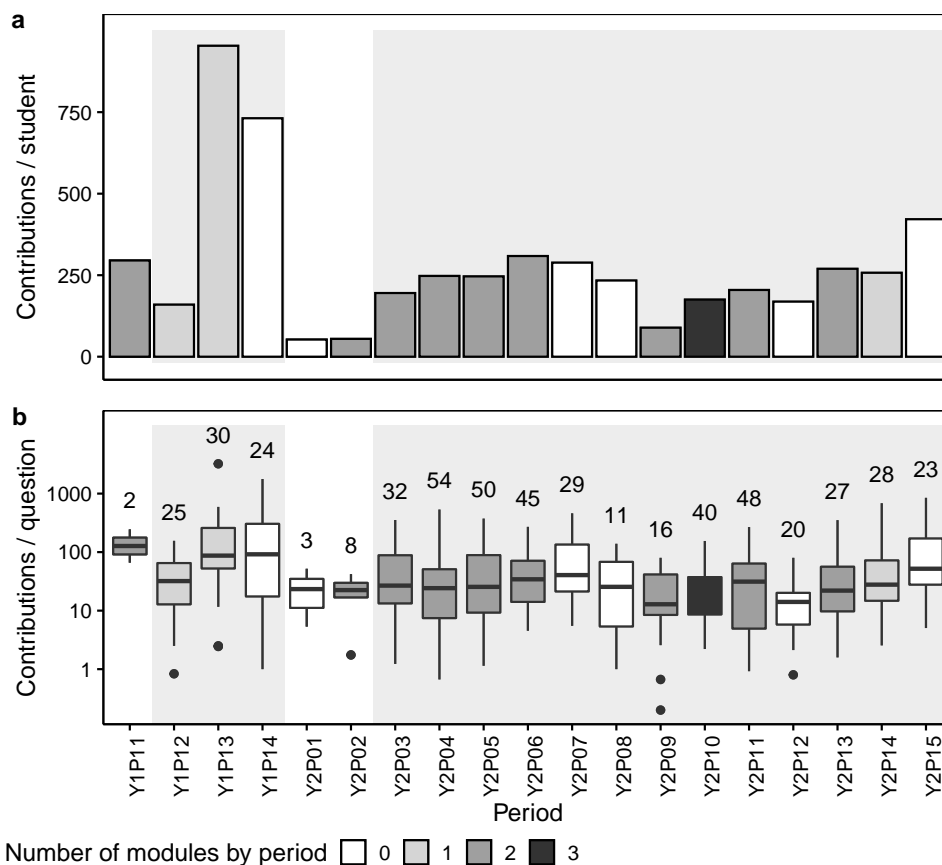


FIGURE 4. a. Average student contributions to the projects per two-week course. b. Contributions per question asked (log scale) for each student as a measurement of the intensity of teacher-learner interactions relative to the progression. Light gray background indicates periods where distance teaching was mandatory due to the COVID-19 lockdown. The number of students that interacted during each period is indicated on top of the boxplots (Y1 is 2019-2020, Y2 is 2020-2021).

progressively, the courses were stretched out over a very long period of time: five successive terms spanning three consecutive years (undergraduate and graduate). That way, the different concepts they had to learn were broken down into subunits (26 modules) that lasted for two weeks each. We also used flipped classrooms and blended teaching and learning (following Spadafora & Zopito’s definition of “any educational model where online delivery ranges from 50% to 80%” [38]), with an emphasis on proactive exchanges with the teachers: students had to ask questions to progress. Overall, these appear to be winning choices because most of our students were successful, excluding a few who failed. Compeau also obtained good results using the flipped classroom with the course targeting students in biology [14].

Despite our course framework, students are more used to a traditional face-to-face approach made of lectures followed by exercises where important concepts are



repeated at the beginning of the practical sessions. They tend to have a passive attitude during lectures and they expect teachers and assistants to feed them with the key concepts. That attitude does not purposely work here. Proactive behavior and development of autonomy are required [18]. They thus have to engage themselves in a very different way of learning. The transition between the theory they read in a book and the projects where they have to apply these concepts is too sharp without a progression that facilitates students' engagement. The four stages of the progression were: (L1) auto-evaluation exercises directly in the online book, (L2) recall of the main concepts and a guided step by step analysis of a first dataset with the learnr tutorials, (L3) at least one guided individual project with another dataset, before (L4) where they are presented yet a different dataset to analyze with limited instructions this time.

#### 4.1. Measured and Perceived Cognitive Workload in learnr Tutorials.

The learnr tutorial was immediately spotted as a key activity in the learning process during the 2018-2019 academic year. As such, we focused our attention on these learnr documents. In 2020-2021, the use of a heuristic engine {gradethis} to provide contextual feedback on the errors students made in their answers was much appreciated.

The number of trials needed to find the right answer is considered here as a workable proxy for the cognitive workload. Contextual feedback allowed the student to correct these answers on their own to find the expected answer. A high number of trials may definitely indicate that students had a hard time to find the right answer. They could misunderstand the concepts, but they could also be trying many solutions at random in order to find the right solution. This indicator alone is not sufficient, so we combined it with the perceived cognitive workload. In the future, the measured RTLX index will serve as a reference to gauge possible optimization of the tutorials, with lower perceived workload without sacrificing the content. The significant decrease in the RTLX value from Course A to C indicates that there is still a margin of progression. We would like to observe such a decrease sooner, perhaps already in the second course. Monitoring the perceived and measured cognitive workload is indeed important “to maintain reflective and systematic approaches in both the development and evaluation of [our] blended approach” [38].

**4.2. Student Activity Profiles with Ongoing Assessment.** Activity tracking in the exercises, primarily set up for the ongoing assessment, also offers the opportunity to study the way learning happens (or not). Indeed, learning analytics provides opportunities to monitor learning events as well as to adjust teaching to improve student outcomes [29, 33], even if they are primarily used to early predict success or failure. In our courses, the failure rate was already rather low and essentially limited to a few defeating students that did not work at all. We were more interested in classifying our participating students according to their behaviors. This paved the way toward a more inclusive pedagogy by spotting different kinds of suboptimal patterns (for instance, never asking questions, looking at hints too quickly without really trying to figure out the answer, being shy about discussing problems on public channels, etc.). Once these patterns are evidenced, we can then consider countermeasures. As Martin and Ndoye highlight from other studies, “benefits that the online learning platform provides with respect to assessment include better monitoring opportunities for student learning and immediate feedback [...], and individual practice opportunities” [29]. As an example, for students that rarely

post their questions publicly, we will test an alternate discussion channel where teachers never post but have read access. In case of an error, the teacher will contact the student privately to explain to him what is wrong. That student would then have the responsibility to reexplain the point correctly to their classmates. This way, this error is never publicly pointed out by the teacher. With tools like the self-organizing map, we should be able to predict suboptimal student profiles early. We could engage in a discussion with the students concerned to determine the cause and find a solution as early as possible. Learning analytics used in this way could promote a differential pedagogical approach, a key for more inclusive teaching [36].

Group projects are one of the keys to our method. Sometimes, groups do not work well, and one student has to do most of the work. This is a clear weakness of this approach, especially if one of the failing students in Fig. 3 Cell (3) is involved. If we could identify the profile of the different students relatively early during the course, we would be able to create better groupings with a blend of different complimentary profiles to enrich the experience of all learners. Maybe should we work exclusively with groups of four to mitigate the impact of one failing student? The balance of heterogeneous and complimentary competences are essential in such a group in order to create mutual emulation and efficacy [31]. Working on group composition will thus be one of our future challenges.

**4.3. Transition from Face-to-face to Distance Learning Imposed by the COVID-19 Lockdown.** Forced distance learning, due to the COVID-19 lockdown did not appear to be a barrier in the production of our students in their projects. A pattern was observed during the first lockdown with a marked decrease in their contributions, followed by a large, compensatory activity. All this happened in a time frame of a couple of weeks. That was the time needed to adapt to the new situation. Several reasons can be hypothesized to explain the adaptation period. Among them, the access to a powerful-enough computer for roughly 15 to 20% of our students during lockdown. In a normal situation, these students had access to computers on the university premises, both in-class and outside of class time. When the lockdown was imposed, these students suffered a lack of hardware. However, to reduce the social numeric divide, the university quickly reacted and computers were lent to them. During the second lockdown, a larger part of our students had acquired their own computer, and solutions were immediately available for the others.

While the contributions/questions remained globally at a similar level in face-to-face and distance learning, their impact on the teacher's timetables was very different. In distance learning, students worked at very different times. Their questions were thus less concentrated during the course periods. Also, an alternation between asynchronous work at home and synchronous work in the computer lab was more beneficial to interactions between students. The social and human components of teaching and learning are key factors that tend to vanish in exclusive distance learning. Contacts through videoconferences only partly compensate for a lack of interactions because in-class presence remains different to video chats. Blended learning combines the best of the two practices if pedagogical setups are accurate and well balanced [7].

**5. Conclusion and Perspectives.** Teaching data science comes with challenges. The discipline is quite young, and we are still seeking the best pedagogical approach.

After three years of teaching data science to undergraduate and graduate students on a biology curriculum with revised pedagogical practices, we have had our first cohort that passed all three courses. There are still two optional courses available in the second year of the Master if they want to push their data science skills further on. However, the three mandatory courses are designed to be self-supported. Globally, most students acquired the expected competences during these courses. We have the feeling that they are more mature and that they effectively acquired the intended outcomes in data science more than with our previous courses in biostatistics, which was given in a more traditional way. The impact of the revised approach to teach biological data science on the way learners manage data and data analysis will be observable during the following years. We will monitor how these students apply their skills in their Master's thesis, and later, in their career or during their PhD thesis. Meanwhile, we will continue to improve our courses by further exploiting the data we accumulate on the activity of our students. Experience gathered during forced distance learning during the COVID-19 lockdown will also be used to improve our courses framework. The radical changes that were required in that context showed that students can accommodate, to a large extent, but also that the diversification of the activities is beneficial to guarantee their engagement [38, 45]. Regarding diversification, in 2020-2021 we successfully tested a kaggle-like challenge (<https://www.kaggle.com/competitions>) in one of the machine learning modules. Such playful activities could also contribute to the diversification of pedagogical practices, interest and motivation of the students [1]. We would also be happy to share experiences with other teachers in data science. Altogether, we are on the way to reshaping the post-COVID teaching landscape, and it will probably be quite different than what we were used to!

## REFERENCES

- [1] C. Alonso-Fernandez, A. Calvo-Morata, M. Freire, I. Martínez-Ortiz and B. Fernandez-Manjon, Applications of data science to game learning analytics data: A systematic literature review, *Computers & Education*, **141** (2019), 103612.
- [2] H. Alvarenga da Silva and A. Sampaio, Teaching introductory statistical classes in medical schools using RStudio and R statistical language: Evaluating technology acceptance and change in attitude toward statistics, *Journal of Statistics Education*, **28** (2020), 212–219.
- [3] L. A. Auken and E. L. Barthelmess, Teaching R in the undergraduate ecology classroom: approaches, lessons learned, and recommendations, *Ecosphere*, **11** (2020), e03060.
- [4] M. Baker, 1,500 scientists lift the lid on reproducibility, *Nature*, **533** (2016), 452–454.
- [5] G. C. Banks, J. G. Field, F. L. Oswald, E. H. O'Boyle, R. S. Landis, D. E. Rupp, S. G. Rogelberg. Answers to 18 Questions About Open Science Practices, *Journal of Business and Psychology*, **34** (2019), 257–270.
- [6] B. Baumer, M. Cetinkaya-Rundel, A. Bray, L. Loi and N. J. Horton, R Markdown: Integrating A Reproducible Analysis Tool into Introductory Statistics, *Technology Innovations in Statistics Education*, **8** (2014).
- [7] R. M. Bernard, E. Borokhovski, R. F. Schmid, R. M. Tamin and P. C. Abrami, A meta-analysis of blended learning and technology use in higher education: from the general to the applied, *J Comput High Educ*, **26** (2014), 87–122.
- [8] C. Boettiger, An Introduction to Docker for Reproducible Research, *SIGOPS Oper. Syst. Rev.*, **49** (2015), 71–79.
- [9] R. Burton, S. Borruat, B. Charlier, N. Coltice, N. Deschryver, F. Docq and al., Vers une typologie des dispositifs hybrides de formation en enseignement supérieur, *Distances et savoirs*, **9** (2011), 69–96.
- [10] J.C. Byers, A. Bittner and S. Hill, Traditional and raw task load index (TLX) correlations: Are paired comparisons necessary?, *Advances in Industrial Ergonomics and Safety*, 1989, 481–485.

- [11] M. Cetinkaya-Rundel and C. Rundel, Infrastructure and Tools for Teaching Computing Throughout the Statistical Curriculum, *American Statistician*, **72** (2018), 58–65.
- [12] M. Cetinkaya-Rundel and V. Ellison, A Fresh Look at Introductory Data Science, *Journal of Statistics Education*, **0** (2021), 1–27.
- [13] W.S. Cleveland, Data science: An action plan for expanding the technical areas of the field of statistics, *International Statistical Review*, **1** (2001), 21–26.
- [14] P. Compeau, Establishing a computational biology flipped classroom, *PLoS Computational Biology*, **15** (2019), 1–8.
- [15] D. Donoho, 50 Years of Data Science, *Journal of Computational and Graphical Statistics*, **26** (2017), 745–766.
- [16] R.A. Estrellado, E. A. Emily, J. Mostipak, J. M. Rosenberg and I. C. Velasquez, *Data science in education using R*, 1st edition, Routledge, London, England, 2020.
- [17] J. Fiksel, L. R. Jager, J. S. Hardin and M. A. Taub Using GitHub Classroom to teach statistics, *Journal of Statistics Education*, **27** (2019), 110–119.
- [18] D. R. Krathwohl, Active learning increases student performance in science, engineering, and mathematics, *Proceedings of the National Academy of Sciences*, **111** (2014), 8410–8415.
- [19] P. Grosjean and G. Engels, Biological data science courses at UMONS, Belgium: student’s activity for 2020-2021 (1.0.0) [Data set], *Zenodo*, (2022).
- [20] P. Grosjean and G. Engels, Biological data science courses at UMONS, Belgium: student’s activity for 2019-2020 (1.0.0) [Data set], *Zenodo*, (2022).
- [21] L.M. Guzman, M.W. Pennell, E. Nikelski and D.S. Srivastava, Successful integration of data science in undergraduate biostatistics courses using cognitive load theory, *CBE-Life Sciences Education*, **18** (2019), ar49 1–10.
- [22] S. G. Hart and L. E. Staveland, Development of NASA-TLX (Task Load Index): Results of Empirical and Theoretical Research, *Advances in Psychology*, **52** (1998), 139–183.
- [23] C. Hsing and V. Gennarelli, Using GitHub in the Classroom Predicts Student Learning Outcomes and Classroom Experiences: Findings from a Survey of Students and Teachers, In *Proceedings of the 50th ACM Technical Symposium on Computer Science Education (SIGCSE ’19)*, 2019, 672–678.
- [24] P.A. Kirschner and C. Hendrick, *How learning happens: seminal works in educational psychology and what they mean in practice*, 1st edition, Routledge, London, New York, 2020.
- [25] T. Kohonen, *Self-Organizing Maps*, 1st edition, Springer-Verlag, Berlin Heidelberg, 1995.
- [26] D. R. Krathwohl, A Revision of Bloom’s Taxonomy: An overview, *Theory Into Practice*, **41** (2002), 212–218.
- [27] K. Larwin and D. Larwin, A meta-analysis examining the impact of computer-assisted instruction on postsecondary statistics education: 40 years of research, *Journal of Research on Technology in Education*, **43** (2011), 253–278.
- [28] L. Lloyd-Smith, Exploring the Advantages of Blended Instruction at Community Colleges and Technical Schools, *Journal of Online Learning and Teaching*, **6** (2010).
- [29] F. Martin and A. Ndoye, Using Learning Analytics to Assess Student Learning in Online Courses, *Journal of University Teaching & Learning Practice*, **13** (2016), 69–96.
- [30] V. Marx, The big challenges of big data, *Nature*, **498** (2013), 255–260.
- [31] R. Mucchielli, *Le travail en equipe*, 1st edition, ESF Editions, Paris, France, 1996.
- [32] J. A. Onwuegbuzie and V. A. Wilson, Statistics Anxiety: Nature, etiology, antecedents, effects, and treatments—A comprehensive review of the literature, *Teaching in Higher Education*, **8** (2003), 195–209.
- [33] C. Romero and S. Vetura, Educational data mining and learning analytics: An updated survey, *WIREs Data Mining Knowl Discov*, **10** (2020), e1355.
- [34] R Core Team (2021), R: A Language and Environment for Statistical Computing, R Foundation for Statistical Computing, Vienna, Austria. <https://www.R-project.org>
- [35] RStudio team (2015), RStudio Cloud, Boston, MA: RStudio, Inc.
- [36] G. Siemens, Learning Analytics: The Emergence of a Discipline, *American Behavioral Scientist*, **57** (2013), 1380–1400.
- [37] B. Sousa and D. Gomes, Teaching With R — A Curse or a Blessing?, In *Proceedings of the Tenth International Conference on Teaching Statistics (ICOTS10 ’18)*, 2018, 672–678.
- [38] N. Spadafora and Z. Marini Self-Regulation and “Time Off”: Evaluations and Reflections on the Development of a Blended Course, *The Canadian Journal for the Scholarship of Teaching and Learning*, **9** (2018).

- [39] A. S. Theobald, A. Hancock and S. Mannheimer, Designing Data Science Workshops for Data-Intensive Environmental Science Research, *Journal of Statistics and Data Science Education*, **29** (2021), S83–S94.
- [40] University of Mons Study programme 2021-2022 : Data Sciences I visualisation and inference, Accessed March 21, 2022, <https://applications.umons.ac.be/web/en/pde/2020-2021/ue/US-B2-SCBIOL-006-M.htm>
- [41] University of Mons Study programme 2021-2022 : Data Sciences II analysis and modelling, Accessed March 21, 2022, <https://applications.umons.ac.be/web/en/pde/2020-2021/ue/US-B3-SCBIOL-006-M.htm>
- [42] University of Mons SStudy programme 2021-2022 : Data Sciences III exploration and prediction, Accessed March 21, 2022, <https://applications.umons.ac.be/web/en/pde/2020-2021/ue/US-M1-BIOECO-004-M.htm>
- [43] R. Wehrens and J. Kruisselbrink, Flexible Self-Organizing Maps in kohonen 3.0, *Journal of Statistical Software, Articles*, **87** (2018), 1–18.
- [44] H. Wickham et al, Welcome to the Tidyverse, *Journal of Open Source Software*, **43** (2019), 1686.
- [45] J. Young, Hybrid teaching seeks to end divide between traditional and online instruction, *The Chronicle of Higher Education*, **48** (2002), A33–34.

Received xxxx 20xx; revised xxxx 20xx.